

PA2 Part 3: Supplemental Material

Update

■ 2024-05-03

- Added slides on how to use `tcsetpgrp` with `setpgid`
- Added slide on an idea on how to implement the parser
- Added slide on how to execute command
- Added slide on how to deal with errors in a pipeline
- Fixed detail in the pipeline slide
- Decreased `JOB_MAX` from 2063244 to 100.
- Added `ARG_MAX`, which is also 100.

Errata

■ Part 1 was updated.

- Made the error message format much easier to understand.
 - » “command: ERROR MESSAGE FROM strerror()” ->
 - » “command: %s” where the string is the return value of strerror().
- Added `_GNU_SOURCE` part to `basename()`.
- Removed the duplicate slide for `longopts`

■ The “**SHELL_NAME**” in this PA should be “**pa2_shell**”.

- But the executable is still called `pa2`.

Parsing

Parsing

■ Grammar

- The inputs passed to the shell follow this grammar:

```
input ::= pipeline ["&"];
pipeline ::= commands | pipeline "|" commands ;
commands ::= command |
            command redirection_operator pathname |
            command redirection_in_operator pathname redirection_out_operator pathname ;
redirection_operator ::= redirection_in_operator | redirection_out_operator ;
redirection_in_operator ::= "<";
redirection_out_operator ::= ">" | ">>";
command ::= (executable | builtin) {args} ;
args ::= { option | option_with_argument | argument } ;
executable ::= executable_in_path | prefixed_self_implemented_executable | path ;
prefixed_self_implemented_executable ::= "pa2_" , self_implemented_executable ;
self_implemented_executable ::= "head" | "tail" | "cat" | "cp" | "mv" | "rm" ;
builtin ::= "cd" | "pwd" | "exit" | "fg" | "bg" | "jobs";
path ::= ["/" | "../" | "/"] pathname ;
```

- Grammar is in EBNF ([...] = optional, | = or, {...} = ...*)

Parsing

■ Structure

- Based on this grammar, we can create these:
 - » Job (Input)
 - » Foreground/Background
 - » Commands[N = 100]
 - » Command
 - » Command (argv[0])
 - » Arguments[N = 100]
 - » Redirection (or NULL)
 - » Built-in, self-implemented, in PATH, or pathname (., .., /)?
 - » Redirection
 - » <, >, >>
 - » pathname

Parsing

■ Using Job Structure

- There is only one job in one input line.
- You can extend the previous input structure by adding pgid, job status (foreground, background, stopped), etc.
- The maximum number of processes in a job is 100 (number from [getrlimit\(2\)](#))
- You can keep track of all the job structures, which can be used for `jobs`.
 - » In bash, this structure is stored using a linked list, but you can use an array since the limit is 8192 and it is easy to manage job numbers with this
 - » Do note that foreground job does not have a job number

Parsing

■ Implementing parsing

- You can use the lex algorithm in *week8* to convert the input into tokens.
- You do not have to use actual parsing algorithms/techniques, especially if you are not familiar with them.
 - » You can just go through each token since the grammar is simple.
 - » But you can still use RD/packrat if you are familiar with them.
- There are two main structures (+ redirection, which can be merged with the command structure)
 - » Job: Has an array with the capacity of 100 commands
 - » Command: Has an array with the capacity of 100 arguments
- So you only have to keep track of two indices: one for command and one for args
 - » And depending on the token, you can change the attribute of the structure
 - » i.e. If & exists at the end, the input is a background job.
 - » i.e., if parser encounters '<', you can set the redirection of the current command to <.

Executing

- **Executing the parsed expression.**
 - Now that you parsed the input, you can easily determine if pipe and redirection are needed.
 - You can now easily determine if the command is a built-in shell command, self-implemented executable, an executable in a given path, or an executable in \$PATH (please note the difference between the last two!), so you can handle a command differently based on its type (if execvp is needed, path has to be modified, etc.)!
 - Assuming your pa2 is in ./bin, your self-implemented variables should be in ./bin, so you should call ./bin/pa2_EXECUTABLE, when pa2_EXECUTABLE is passed as a command.

Pipeline

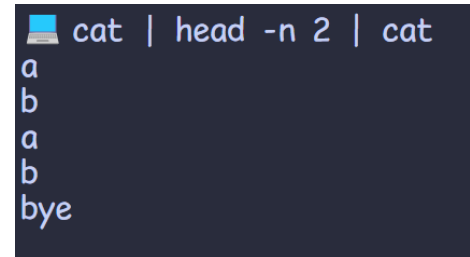
Pipeline

■ Pipeline

- `command1 | command2 | command3`
 - » `command1`'s input is STDIN (of shell or redirected STDIN)
 - » `command1`'s output is `command2`'s input
 - » `command2`'s output is `command3`'s input
 - » `Command3`'s output is STDOUT (of shell or redirected STDOUT)
- All of these commands are running concurrently.

» If we have “`cat | head -n 2 | cat`”, what will happen?

- » First line: `a` -> goes to first `cat`
- » Second line: `b` -> goes to second `cat`
- » `Head` reads two lines and can safely quit
- » Second `cat` prints the two line and exits
- » First `cat` is still running since it is expecting input.
- » After pressing enter, it finds out pipe has been closed and exits.



```
cat | head -n 2 | cat
a
b
a
b
bye
```

Pipeline

■ Pipe

- Short counts are very normal in pipes as seen in the previous example.
 - » When data is available in STDIN, programs will usually process that and then read STDIN again if applicable.
- If we have ``a | b``,
 - » What happens if ``a`` terminates while ``b`` is waiting for input?
 - » ``b`` will receive an EOF in the pipe and then terminate as STDIN has essentially been closed.
 - » What happens if ``b`` terminates and `a` attempts to write to the pipe?
 - » ``a`` will expect an input (read function blocks!), attempt to write the input, but receive SIGPIPE as the pipe has been closed already.
 - » Example: ``cat | exit``
 - » Although ``cat`` exits last, the exit code is the exit code of ``exit``, not ``cat``

Pipeline

▪ Handling errors

- What happens if an error occur?
 - » The process will exit and usually print an error to stderr before exiting.
 - » The same process happens as the previous slide (i.e., if `a` terminates, then b terminates as well, but if `b` terminates, then `a` will still expect input)
 - » If the command does not exist, it should still print “COMMAND: command not found”

tcsetpgrp

tcsetpgrp – basic example

```
int main() {
    pid_t pid;
    int shell_fd = STDIN_FILENO;

    switch (pid = fork()) {
        case -1:
            exit(1);
        case 0: {
            pid_t pgid = setsid();
            tcsetpgrp(shell_fd, pgid);
            execlp("ls", "ls", NULL);
            exit(0);
            break;
        }
        default: {
            pid_t child_pid = pid;
            pid_t child_pgid = getpgid(child_pid);
            pid = getpid();
            pid_t pgid = getpgid(pid);
            tcsetpgrp(shell_fd, child_pgid);
            waitpid(child_pid, NULL, 0);
            tcsetpgrp(shell_fd, pgid);
            puts("Works well!");
        }
    }
    exit(0);
    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/wait.h>
#include <unistd.h>
```

- tcsetpgrp should be called in both parent and child
- Process should call tcsetpgrp when it wants to control the terminal
- Note: execlp is used here instead of execvp to reduce lines. You can use execvp.

tcsetpgrp – issue

```
int main() {
    pid_t pid;
    int shell_fd = STDIN_FILENO;

    switch (pid = fork()) {
        case -1:
            exit(1);
        case 0: {
            pid_t pgid = setsid();
            tcsetpgrp(shell_fd, pgid);
            execlp("cat", "cat", NULL);
            exit(0);
            break;
        }
        default: {
            pid_t child_pid = pid;
            pid_t child_pgid = getpgid(child_pid);
            pid = getpid();
            pid_t pgid = getpgid(pid);
            tcsetpgrp(shell_fd, child_pgid);
            kill(child_pid, SIGTSTP);
            waitpid(child_pid, NULL, 0);
            tcsetpgrp(shell_fd, pgid);
            puts("Works well!");
        }
    }
    exit(0);

    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/wait.h>
#include <unistd.h>
```

- This SIGTSTP signal will cause the process not to terminate.
- “Works well” will not be printed.

tcsetpgrp

■ Issue

- This example does not work well with SIGTSTP as waitpid does not wait for stopped processes
- Solution?
 - » WUNTRACED flag in waitpid
 - » Stopped processes are considered as well.
 - » How to differentiate regular termination, SIGINT, SIGTSTP, and other signals? -> Use status from waitpid
 - » WIFEXITED(status): normal termination
 - » WIFSIGNALED(status): terminated due to signal
 - » WTERMSIG(status): returns signal number that caused the process to terminate
 - » WIFSTOPPED(status): stopped (not terminated)
 - » WSTOPSIG(status): returns signal number that caused the process to stop

tcsetpgrp – solution

```
int main() {
    pid_t pid;
    int shell_fd = STDIN_FILENO;

    switch (pid = fork()) {
        case -1:
            exit(1);
        case 0: {
            pid_t pgid = setsid();
            tcsetpgrp(shell_fd, pgid);
            execlp("cat", "cat", NULL);
            exit(0);
            break;
        }
        default: {
            pid_t child_pid = pid;
            pid_t child_pgid = getpgid(child_pid);
            pid = getpid();
            pid_t pgid = getpgid(pid);
            tcsetpgrp(shell_fd, child_pgid);
            kill(child_pid, SIGTSTP);
            waitpid(child_pid, NULL, WUNTRACED);

            tcsetpgrp(shell_fd, pgid);
            puts("Works well!");
        }
    }
    exit(0);

    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/wait.h>
#include <unistd.h>
```

- By adding WUNTRACED, “Works well” can now be printed.
- Now that this job has been stopped, we can set the status of this job to “Stopped” and we can move this job to the background or foreground.

tcsetpgrp

■ Moving job back to foreground

- Shell has to let go of the terminal and give the control to this job
- After doing this, you have to pass SIGCONT to the progress group of the job
 - » Why progress group?
 - » So that all the subprocesses in the job (pipeline) can continue
- And then the shell should wait for the process to finish and then retrieve the terminal afterwards

tcsetpgrp – continue

```
int main() {
    pid_t pid;
    int shell_fd = STDIN_FILENO;

    switch (pid = fork()) {
        case -1:
            exit(1);
        case 0: {
            pid_t pgid = setsid();
            tcsetpgrp(shell_fd, pgid);
            execlp("cat", "cat", NULL);
            exit(0);
            break;
        }
        default: {
            pid_t child_pid = pid;
            pid_t child_pgid = getpgid(child_pid);
            pid = getpid();
            pid_t pgid = getpgid(pid);
            tcsetpgrp(shell_fd, child_pgid);
            kill(child_pid, SIGTSTP);
            waitpid(child_pid, NULL, WUNTRACED);
            tcsetpgrp(shell_fd, pgid);
            puts("Works well!");
            // Continue job
            tcsetpgrp(shell_fd, child_pgid);
            kill(child_pid, SIGCONT);
            waitpid(child_pid, NULL, WUNTRACED);
            tcsetpgrp(shell_fd, pgid);
        }
    }
    exit(0);

    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/wait.h>
#include <unistd.h>
```

tcsetpgrp

▪ However, this creates a new session

- We need to keep the process groups in the same session!
- So we have to use `setpgid(pid, pgid)` instead of `setsid()`
 - » `pid = child pid`
 - » `pgid = pid` (if new job) or `pgid` (if stopped job)
 - » So somehow store the `pid` and `pgid` in a structure
 - » You have to call this in both the child and parent processes
- But this introduces a new problem!

If `tcsetpgrp()` is called by a member of a background process group in its session, and the calling process is not blocking or ignoring `SIGTTOU`, a `SIGTTOU` signal is sent to all members of this background process group.

tcsetpgrp – issue

```
int main() {
    pid_t pid;
    int shell_fd = STDIN_FILENO;
    switch (pid = fork()) {
        case -1:
            exit(1);
        case 0: {
            pid_t pid = getpid();
            setpgid(pid, pid); // or pgid if it exists
            puts("This will stop...");
            tcsetpgrp(shell_fd, pid);
            puts("Now I can call cat! Press Ctrl+D to exit!");
            execlp("cat", "cat", NULL);
            exit(0);
            break;
        }
        default: {
            pid_t child_pid = pid;
            pid_t child_pgid = pid;
            pid = getpid();
            pid_t pgid = getpgid(pid);
            tcsetpgrp(shell_fd, child_pgid);
            kill(child_pid, SIGTSTP);
            waitpid(child_pid, NULL, WUNTRACED);
            tcsetpgrp(shell_fd, pgid);
            // Continue again
            tcsetpgrp(shell_fd, child_pgid);
            puts("Parent was stopped");
            kill(child_pid, SIGCONT);
            waitpid(child_pid, NULL, WUNTRACED);
            puts("Parent will enter a stopped state again...");
            tcsetpgrp(shell_fd, pgid);
            puts("Bye!");
        }
    }
    exit(0);
    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/wait.h>
#include <unistd.h>
```

tcsetpgrp

- **Process will enter a stopped state twice**
 - tcsetpgrp will trigger a SIGTTOU signal if called by a background process group in the same session
 - » Did not happen before as the sessions were different
 - How to handle this?
 - » Ignore SIGTTOU completely in parent
 - » In child,
 - » ignore SIGTTOU before calling tcsetpgrp
 - » BUT, set it to the default behaviour after calling tcsetpgrp
 - » Keep in mind, you have to ignore other signals aside from SIGTTOU too in the parent while keeping the default behaviour in the child!
 - In this way, you can safely use tcsetpgrp with multiple process groups in the same session.