

# **PA2 Part 2: Extending the Shell**

# PA2

## ■ Goal

- Final goal of this assignment is to create a shell with redirection and pipe support.
- For now, we will extend our shell by
  - » Pipe and redirection support
  - » Job control
    - » Using concepts like process groups, sessions, background processes, etc.
    - » Handling signals
  - » Adding other built-in commands (bg, fg, kill, cd, exit, and pwd)
- Hint:  
[https://www.gnu.org/software/libc/manual/html\\_node/Implementing-a-Shell.html](https://www.gnu.org/software/libc/manual/html_node/Implementing-a-Shell.html) (don't copy the code<sup>2</sup>)

# PA2

- **Supplement resources/guides (don't copy the code)**
  - [https://www.gnu.org/software/libc/manual/html\\_node/Implementing-a-Shell.html](https://www.gnu.org/software/libc/manual/html_node/Implementing-a-Shell.html)
  - <https://www.linusakesson.net/programming/tty/>
    - » Focuses on tty, which we will not learn, instead of shell, but explains job control and signals very well.

# Shell



# Shell

## ▪ Shell

- A shell reads user input, evaluates the input, and then return the result.
  - » When evaluating the input, the shell can first lex the input, and then parse the result to easily execute the commands.
- For PA2, you will be implementing a shell that is heavily based on bash.
  - » Thus, you can use bash to test the behaviour of your shell but note that not all features in bash are found in your shell (e.g. variables, expansion, “~”, etc.).

# Shell

## ■ Grammar

- The inputs passed to the shell follow this grammar:

```
input ::= pipeline ["&"];
pipeline ::= commands | pipeline "|" commands ;
commands ::= command |
            command redirection_operator pathname |
            command redirection_in_operator pathname redirection_out_operator pathname ;
redirection_operator ::= redirection_in_operator | redirection_out_operator ;
redirection_in_operator ::= "<";
redirection_out_operator ::= ">" | ">>";
command ::= (executable | builtin) {args} ;
args ::= { option | option_with_argument | argument } ;
executable ::= executable_in_path | prefixed_self_implemented_executable | path ;
prefixed_self_implemented_executable ::= "pa2_", self_implemented_executable ;
self_implemented_executable ::= "head" | "tail" | "cat" | "cp" | "mv" | "rm";
builtin ::= "cd" | "pwd" | "exit" | "fg" | "bg" | "jobs";
path ::= ["/" | "../" | "/"] pathname ;
```

- Grammar is in EBNF ( [...] = optional, | = or, {...} = ...\* )

# Shell

## ■ Input

- User enters input up to 200 bytes.
- This input can contain whitespace, but it is ignored.
- Before being processed, this input can be lexed and parsed.
  - » An example function to lex/tokenize the input can be found in the skeleton code for week 8.

## ■ Unknown command

- If the shell encounters an unknown command, it should print “COMMAND: command not found”

# Shell

## ■ Pipeline

- Set of concurrent processes called subprocesses that are connected through standard streams
  - » `$ command1 | command2 | command3`
    - » `command1`'s output is `command2`'s input
    - » `command2`'s output is `command3`'s input
  - » These processes are independent
    - » ``$ echo 1 | cat | exit | echo 2 | cat`` returns 2 and does not exit the shell. `exit` only exits the subprocess.
    - » Do not run these commands “sequentially”; first create a process for each command and wait for all of them to finish
- Implemented using pipes
- Associated with a job.



# Shell

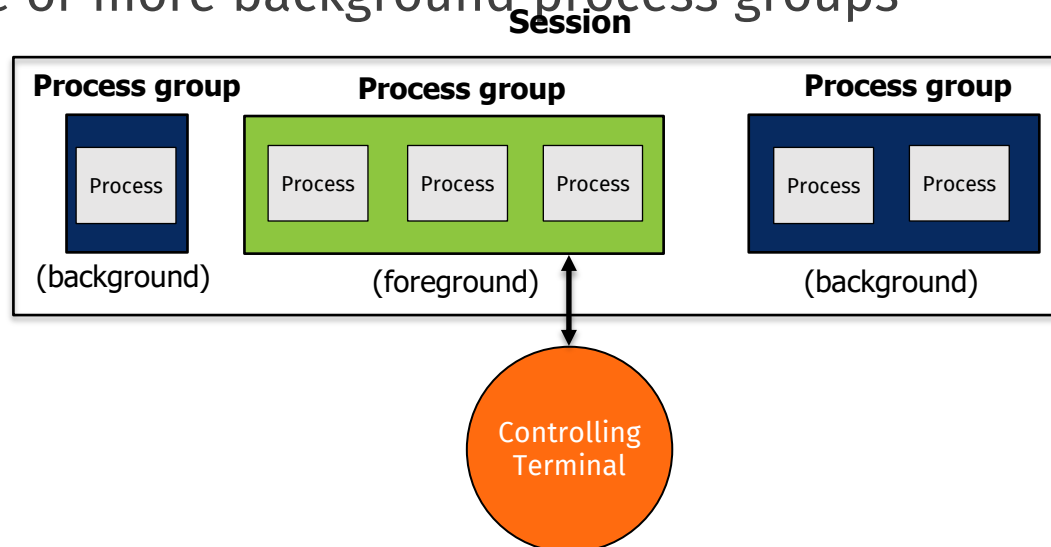
## ▪ Redirection

- Input redirection (<)
  - » Replaces standard input of process with file in [pathname]
- Output redirection (>, >>)
  - » Replaces standard output of process with file in [pathname]
  - » If redirected with >, the file is initialized and rewritten
  - » If redirected with >>, append from the end of file (if it exists)
- If an error occurs due to the file, print “SHELL\_NAME: pathname: strerror(errno)”

# Job Control

# Session (recall)

- **Collection of one or more process groups**
  - Session can have a single terminal
    - » This terminal is called a controlling terminal
  - Process group within a session can be divided into:
    - » A single foreground process group
    - » One or more background process groups



# Jobs

## ■ Terminology

- Job: “A set of processes, comprising a shell pipeline, and any processes descended from it, that are all in the same process group.” - POSIX Issue 7
- Job can be a background group or foreground group.
- Job can have one or more concurrent process(es) via a pipeline (|).
- A job table is used to keep track of all the jobs, which has a limit of 8192.
- Note that this means that shell PGID != Job PGID.

# Jobs

- This is how bash holds processes and jobs in the actual code. (jobs.h). This is just to give you an idea on how it is done in actual applications.

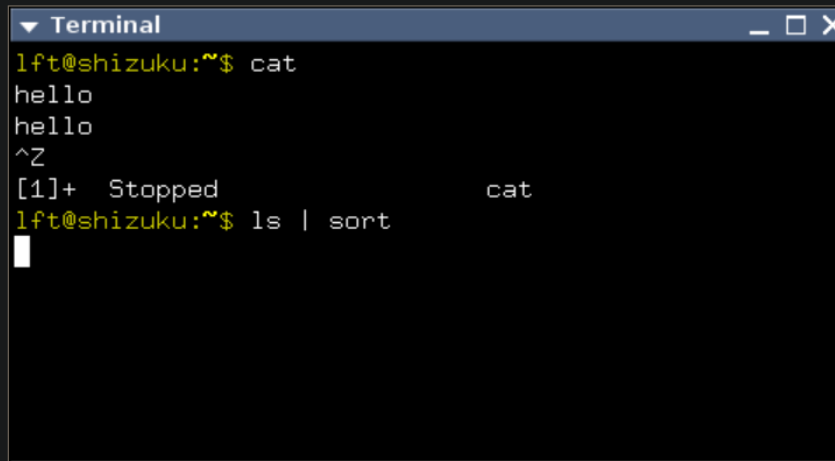
```
typedef struct process {
    struct process *next; /* Next process in the pipeline. A circular chain. */
    pid_t pid; /* Process ID. */
    WAIT status; /* The status of this command as returned by wait. */
    int running; /* Non-zero if this process is running. */
    char *command; /* The particular program that is running. */
} PROCESS;

typedef struct job {
    char *wd; /* The working directory at time of invocation. */
    PROCESS *pipe; /* The pipeline of processes that make up this job. */
    pid_t pgrp; /* The process ID of the process group (necessary). */
    JOB_STATE state; /* The state that this job is in. */
    int flags; /* Flags word: J_NOTIFIED, J_FOREGROUND, or J_JOBCONTROL. */
#ifdef JOB_CONTROL
    COMMAND *deferred; /* Commands that will execute when this job is done. */
    sh_vptrfunc_t *j_cleanup; /* Cleanup function to call when job marked JDEAD */
    PTR_T cleanarg; /* Argument passed to (*j_cleanup)() */
#endif /* JOB_CONTROL */
} JOB;
```

# Jobs

- Taken from “The TTY demystified”

The following shell interactions...



```
lft@shizuku:~$ cat
hello
hello
^Z
[1]+  Stopped                  cat
lft@shizuku:~$ ls | sort
```

...correspond to these processes...

Session 100		Session 101			
Job 100		Job 101	Job 102	Job 103	
XTerm (100)		bash (101)	cat (102)	ls (103)	sort (104)
stdin: -		stdin: /dev/pts/0	stdin: /dev/pts/0	stdin: /dev/pts/0	stdin: pipe0
stdout: -		stdout: /dev/pts/0	stdout: /dev/pts/0	stdout: pipe0	stdout: /dev/pts/0
stderr: -		stderr: /dev/pts/0	stderr: /dev/pts/0	stderr: /dev/pts/0	stderr: /dev/pts/0
PPID: ?		PPID: 100	PPID: 101	PPID: 101	PPID: 101
PGID: 100		PGID: 101	PGID: 102	PGID: 103	PGID: 103
SID: 100		SID: 101	SID: 101	SID: 101	SID: 101
TTY: -		TTY: /dev/pts/0	TTY: /dev/pts/0	TTY: /dev/pts/0	TTY: /dev/pts/0

# Foreground processes

## ▪ Synchronous commands

- Normal jobs that run in the foreground and control the terminal -> the shell no longer runs in the foreground.
- `$ grep aladdin Aladdin.txt | sort > sorted.txt`
  - » This single job will run in the foreground.
  - » Each command in the pipeline is referred to as a subprocess.
  - » The first subprocess is the process group leader, but all subcommands are running simultaneously.
- The exit status is the status of the last command.

# Foreground processes

- `int tcsetpgrp(int fd, pid_t pgrp);`
  - Sets the given **pgid** as the foreground process group with the controlling terminal associated with **fd** (i.e. `STDIN_FILENO`)
  - Has to be called in both parent and child processes.
  - The shell has to regain the controlling terminal after the child has been terminated or stopped.
  - Returns **SIGTTOU** signal to the shell, which will stop it, so it has to be ignored by the shell.
  - A bit tricky to use --- an example code will be added soon



# Background processes

## ▪ & (asynchronous commands)

- & can be added to the end of a job to make run in the background (i.e. shell does not wait for the job to finish)
- `$ grep aladdin Aladdin.txt | sort > sorted.txt &`
  - » This single job will run in the background.
- It will print “[JOB ID] LAST\_PROCESS\_PID” (i.e. [1] 100)
- Returns 0 as an exit status

# Signals

## ■ Signals

- Recall that signals are sent to all processes in the process group.
- If a job is currently running, the shell is running in the background, so it will not receive the signal.
- However, it will still receive **SIGCHLD**, which means a child has been stopped or terminated.
  - » This signal can be used to handle (completed/stopped) jobs.
- You can use the macros **WIFEXITED**, **WIFSIGNALED** along with **WTERMSIG**, and **WIFSTOPPED** along with **WSTOPSIG** to determine if the process has terminated without a signal, terminated due to a signal like SIGINT, or stopped due to a signal like SIGTSTP.

# Signals

## ▪ Signals

- A shell should not terminate or stop when these signals are received by the shell:
  - » SIGINT, SIGQUIT, SIGTSTP, SIGTTIN, SIGTTOU.
- However, the child processes should still receive these signals (SIG\_DFL).

# bg

## ■ Synopsis

- `bg [JOB NUMBER...]`

## ■ Description

- “bg sends jobs to the background, resuming them if they are stopped”
- If JOB NUMBER is not passed, the most recent job would be put in the background.
- JOB NUMBER format: %1, %2, %3, etc.

## ■ Error

- If there are no jobs that can be put in the background, print “SHELL\_NAME: bg: current: no such job”
- If the given job is invalid, print “SHELL\_NAME: bg: JOB NUMBER: no such job” (still run other valid jobs in bg)

# fg

## ■ Synopsis

- fg [JOB NUMBER]

## ■ Description

- fg sends a “job to the foreground, resuming them if it is stopped”
- If JOB NUMBER is not passed, the most recent job would be put in the foreground.
- JOB NUMBER format: %1, %2, %3, etc.

## ■ Error

- If there are no jobs that can be put in the foreground, print “SHELL\_NAME: fg: current: no such job”
- If the given job is invalid, print “SHELL\_NAME: fg: JOB  
NUMBER: no such job”

# jobs

- **Synopsis**

- jobs

- **Description**

- “jobs prints a list of the currently running jobs and their status”
  - Format for each job: [JOB ID]SYMBOL STATE COMMAND
    - » i.e. [1]+ Running sleep 10 &
  - SYMBOL:
    - » + -> Current/most recent job (i.e., if you call fg, this job will be called)
    - » - -> The second-most recent job
    - » NONE -> Other jobs

# pwd

## ■ Synopsis

- `pwd [OPTION]`

## ■ Description

- Outputs the current working directory
- `-L` prints logical directory ( `getenv("PWD")` )
- `-P` prints physical directory ( `getcwd()` )
- Prints the logical directory by default

# pwd

- `char *getcwd(char buf[.size], size_t size);`
  - Gets working directory
  - Return value:
    - » Pointer to string on success. (should be buf)
    - » NULL with errno on failure
  - Size can be set to PATH\_MAX (limits.h), which is usually set to 4096 bytes.



# cd

## ■ Synopsis

- `cd [DIRECTORY]`

## ■ Description

- “cd changes the current working directory”
- If DIRECTORY is not passed, change the directory to the home directory (i.e., `getenv("HOME")`)

## ■ Errors

- When directory does not exist, “SHELLNAME: cd: DIRECTORY: No such file or directory”



- `int chdir(const char *path);`
  - Changes working directory to path.
  - Return value: On success, 0. On error, -1 with errno.
  - Possible Errors:
    - » EACCES: not enough search (execute) permissions for one of the ancestors of path.
    - » ENOENT: directory does not exist
    - » ENOTDIR: ancestor in path is not a directory

# exit

## ■ Synopsis

- `exit [CODE]`

## ■ Description

- exits the shell
- Exit with CODE if given. If CODE  $\geq$  255, exit with 255.
- Otherwise, exit with the exit status of the last command executed (usually 0, but not always).

## ■ Errors

- When code is not a number, ``SHELL_NAME: CODE: numeric argument required`` and exit with status of 2
- When more than one argument is passed, print ``SHELL NAME: exit: too many arguments``